Self-supervised confidence-aware denoising imputation of genomic data

SFU

Mehdi Foroozandeh, Abdul Rahman Diab, and Maxwell Libbrecht

The problem

Most (cell type, assay) experiments were never run, and those that were differ in depth, read length, run type and platform. In the ENCODE Imputation Challenge, a baseline that averages each assay over other cell types beat all but two submitted methods [1].

CANDI

Confidence-**A**ware **N**eural **D**enoising **I**mputer [2] is a self-supervised transformer over raw read counts, DNA sequence and four experimental covariates. It predicts a distribution for every assay in every 25 bp bin, so each imputed value has an interval. It has no cell-type embedding, so it runs on cell types it has never seen.

Key findings

- CANDI is **1st of 26** on the ENCODE Imputation Challenge blind test.
- CANDI's intervals are calibrated, and its latent features are predictive of gene expression.

CANDI schematic workflow

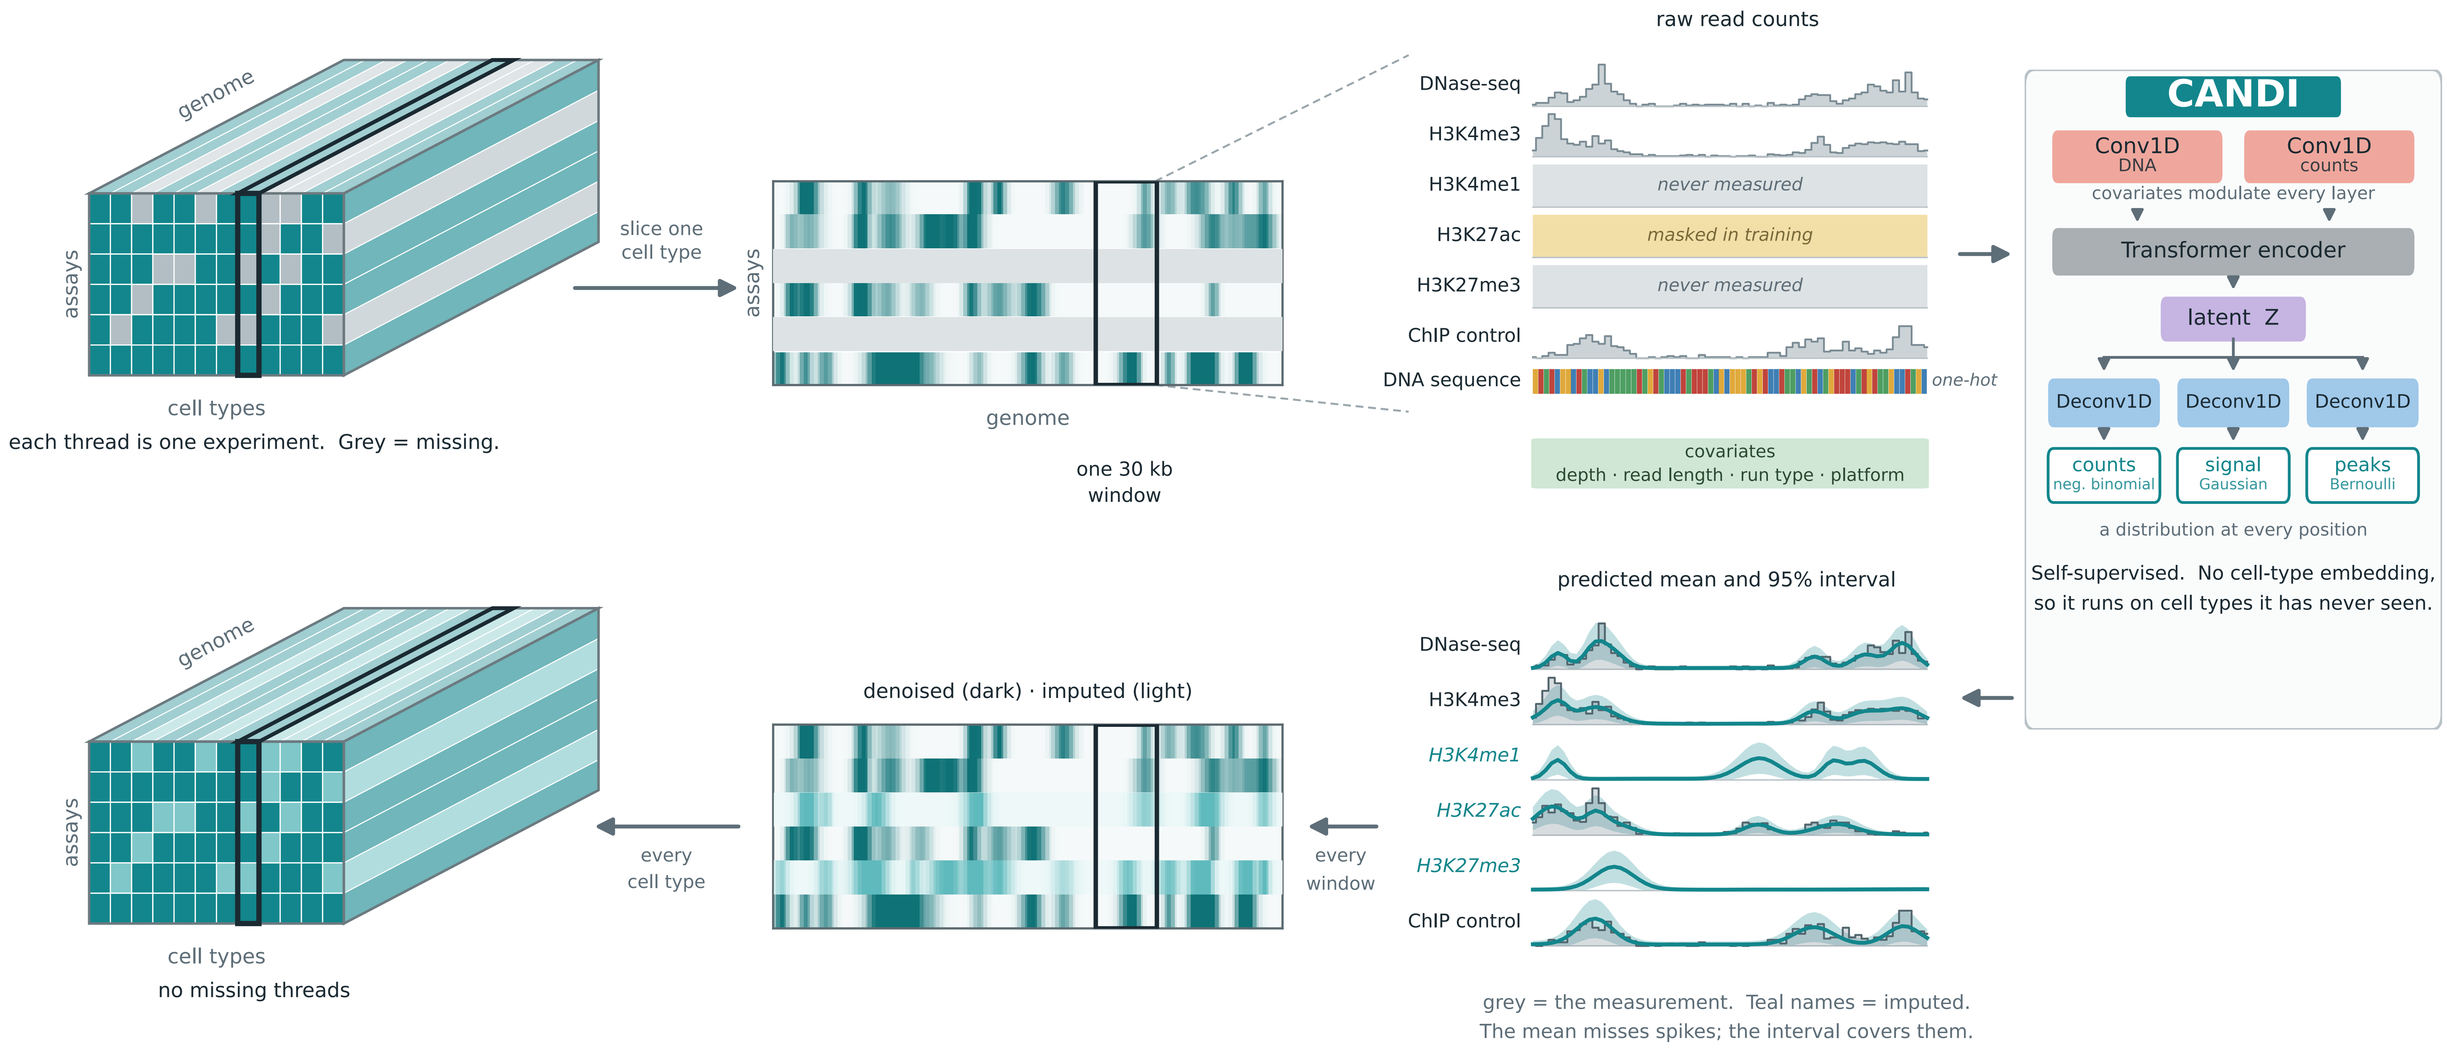

Top: one cell type is sliced out of the cell type × assay × genome tensor, and a 30 kb window of it goes into CANDI as raw counts, DNA and covariates. Grey assays were never measured; yellow ones are masked in training. **Bottom:** CANDI returns a mean and a 95% interval for every assay, and window by window this fills every missing thread of the tensor.

Downstream utility of CANDI

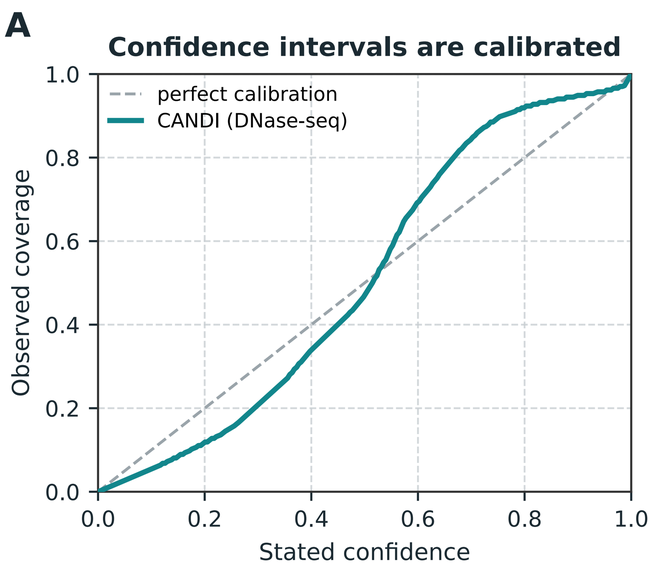

A For DNase-seq (cell line DND-41), the observed coverage of CANDI's intervals follows their stated confidence over the full range.

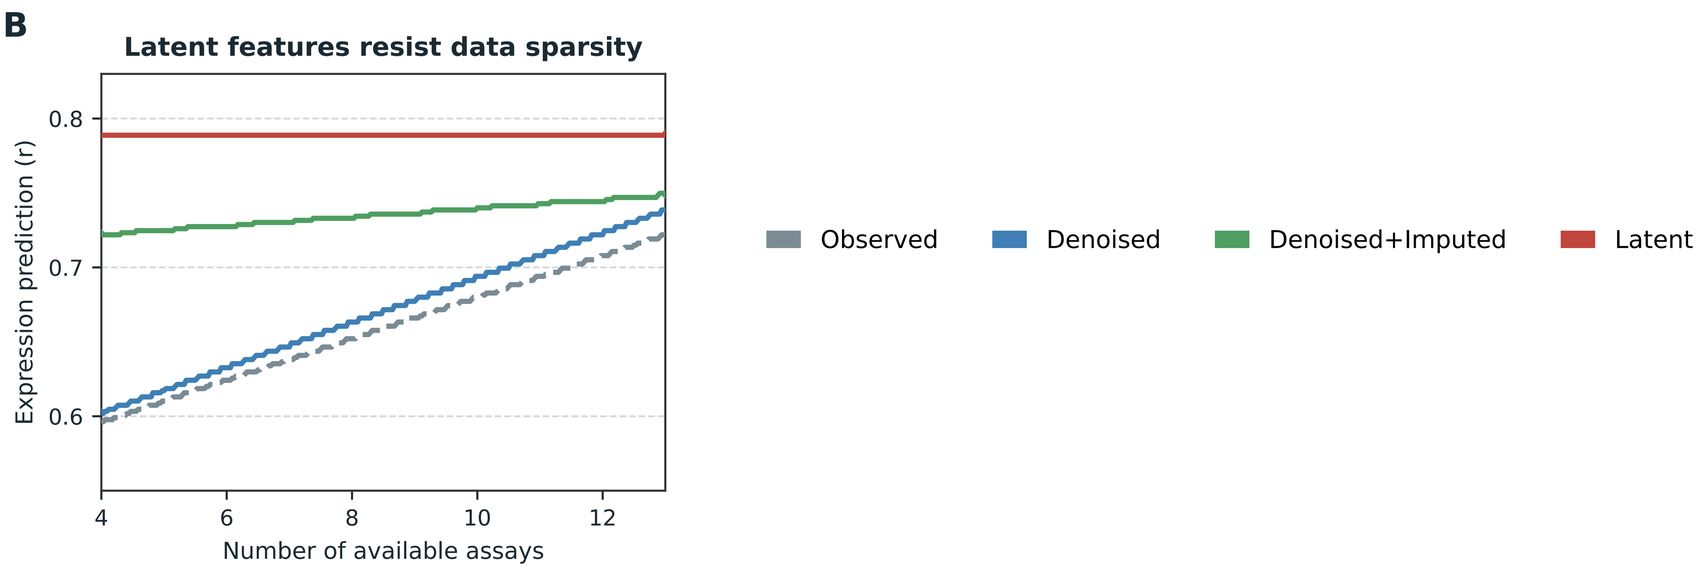

B Gene expression predicted from each feature set as assays are removed. Latent features hold $r \approx 0.79$ from 13 assays down to 4; observed data falls to 0.60.

Fixing the signal-level shift in the ENCODE Imputation Challenge blind test

51 blind-test experiments (12 cell types, 8 assays), scored with the challenge's own nine measures and ranking statistic; our scorer reproduces every published entry to a median relative difference of 9×10^{-10}. CANDI's output as it comes ranks 14th, so for each assay we pass it through one fixed layer, $\max(\beta_0 + \beta_1 \cdot \text{CANDI} + \beta_2 \cdot \text{baseline},\ 0)$, whose 24 numbers are fitted by least squares on the 133 training experiments only; with it, CANDI ranks 1st.

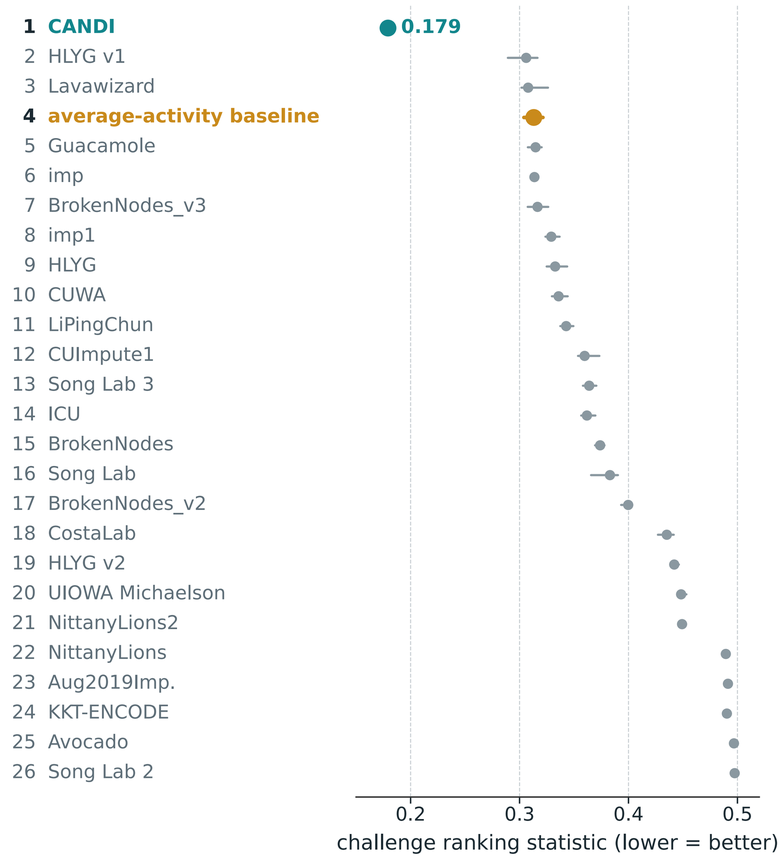

D The challenge's leaderboard with CANDI added. CANDI is first at 0.179 (bootstrap range 0.177–0.184); the next entry scores 0.306. Bars span the ten bootstrap groups.

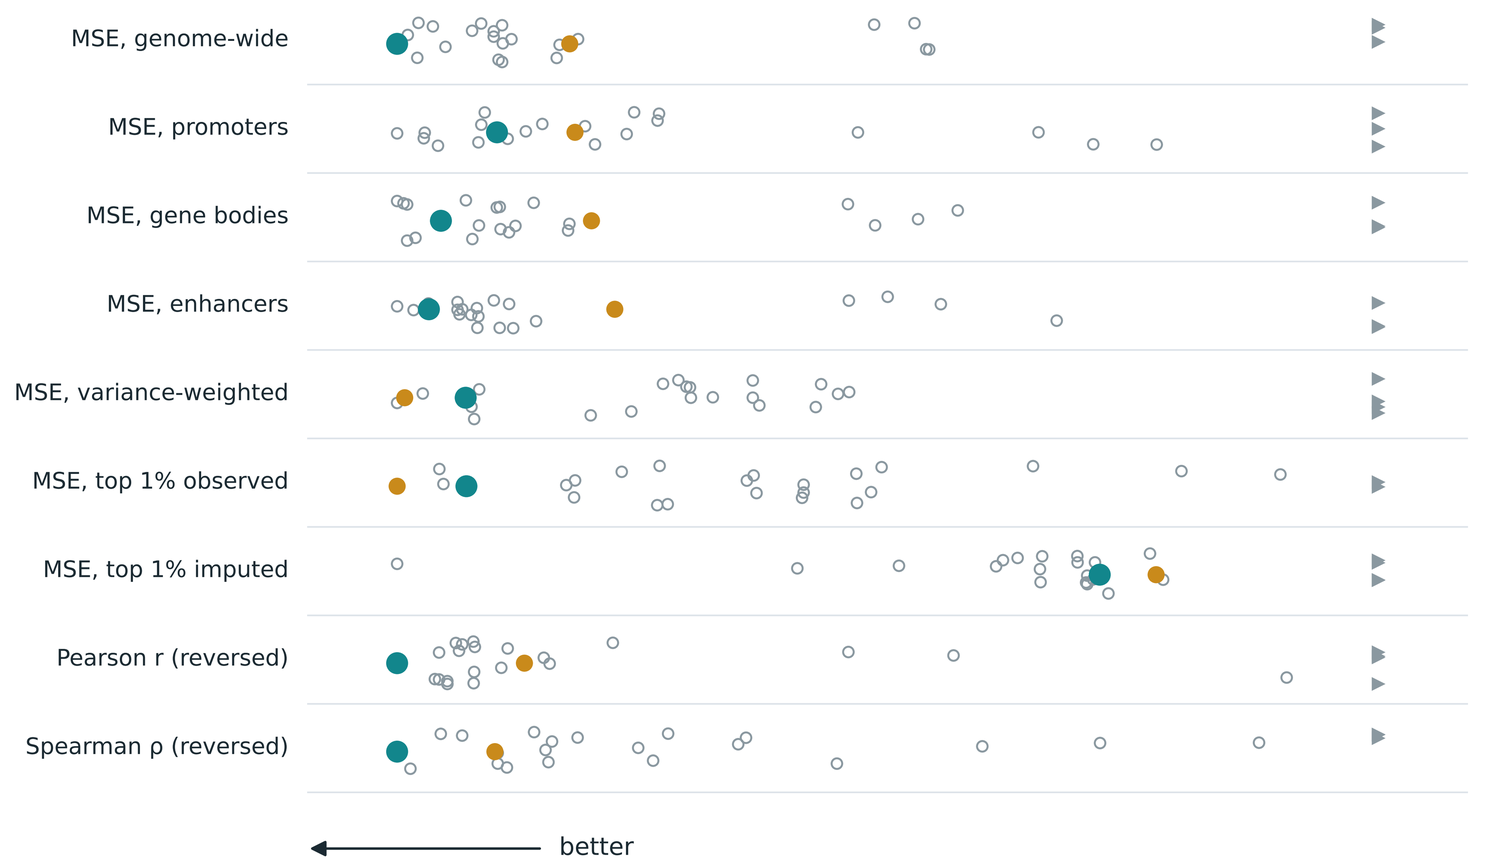

E Each of the nine measures relative to the baseline, median over the 51 experiments. CANDI is first on genome-wide MSE, Pearson and Spearman, and in the top 8 on five more.

mehdiforoozandeh.github.io/CANDI

[1] Schreiber J, Boix C, *et al.* The ENCODE Imputation Challenge: a critical assessment of methods for cross-cell type imputation of epigenomic profiles. *Genome Biol* 24:79 (2023).

[2] Foroozandeh M, Diab AR, Libbrecht M. CANDI: self-supervised, confidence-aware denoising imputation of genomic data. *bioRxiv* 2025.01.23.634626 (2025).

Panels D–E: blind-test scoring from the Libbrecht lab's EIC reproduction (experiments 017, 025). *Baseline* is the challenge's average-activity track. The output correction was chosen after the challenge results were public; none of its numbers is fitted on blind-test data.